

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: .read_csv(r"C:\Users\dinesh reddy\AppData\Local\Microsoft\Windows\InetCache\IE\SYERC185\Mobile_Price_Classification_train[1].csv")
```

Out[2]:

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	talk_time
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	2549	9	7	19
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	2631	17	3	7
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	2603	11	2	9
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	2769	16	8	11
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	1411	8	2	15
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	1890	668	13	4	19
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	1965	2032	11	10	16
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	1632	3057	9	1	5
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	670	869	18	10	19
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	754	3919	19	4	2

2000 rows × 21 columns

```
In [3]: test_df=pd.read_csv(r"C:\Users\dinesh reddy\AppData\Local\Microsoft\Windows\InetCache\IE\3UKEA11U\Mobile_Price_Classification_test\test_df")
```

Out[3]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w	talk_time
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12	7	2
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6	0	7
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17	10	10
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10	0	7
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15	8	7
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14	8	15
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8	1	19
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5	0	14
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15	11	6
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9	2	3

1000 rows × 21 columns

```
In [4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [5]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              1000 non-null   int64
1   battery_power    1000 non-null   int64
2   blue            1000 non-null   int64
3   clock_speed      1000 non-null   float64
4   dual_sim         1000 non-null   int64
5   fc              1000 non-null   int64
6   four_g           1000 non-null   int64
7   int_memory       1000 non-null   int64
8   m_dep            1000 non-null   float64
9   mobile_wt        1000 non-null   int64
10  n_cores          1000 non-null   int64
11  pc               1000 non-null   int64
12  px_height        1000 non-null   int64
13  px_width         1000 non-null   int64
14  ram              1000 non-null   int64
15  sc_h             1000 non-null   int64
16  sc_w             1000 non-null   int64
17  talk_time        1000 non-null   int64
18  three_g          1000 non-null   int64
19  touch_screen     1000 non-null   int64
20  wifi             1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [6]: x=test_df.drop('dual_sim',axis=1)
        y=test_df['dual_sim']
```

```
In [7]: x=test_df.drop('dual_sim',axis=1)
        y=test_df['dual_sim']
```

```
In [8]: train_df['blue'].value_counts()
```

```
Out[8]: blue
0      1010
1       990
Name: count, dtype: int64
```

```
In [9]: test_df['blue'].value_counts()
```

```
Out[9]: blue
1      516
0      484
Name: count, dtype: int64
```

```
In [10]: T={"three_g":{"Yes":1,'No':0}}
train_df=train_df.replace(T)
print(train_df)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	842	0	2.2	0	1	0	7	\
1	1021	1	0.5	1	0	1	53	
2	563	1	0.5	1	2	1	41	
3	615	1	2.5	0	0	0	10	
4	1821	1	1.2	0	13	1	44	
...	
1995	794	1	0.5	1	0	1	2	
1996	1965	1	2.6	1	0	0	39	
1997	1911	0	0.9	1	1	1	36	
1998	1512	0	0.9	0	4	1	46	
1999	510	1	2.0	1	5	1	45	

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	
0	0.6	188	2	...	20	756	2549	9	7	\
1	0.7	136	3	...	905	1988	2631	17	3	
2	0.9	145	5	...	1263	1716	2603	11	2	
3	0.8	131	6	...	1216	1786	2769	16	8	
4	0.6	141	2	...	1208	1212	1411	8	2	
...	
1995	0.8	106	6	...	1222	1890	668	13	4	
1996	0.2	187	4	...	915	1965	2032	11	10	
1997	0.7	108	8	...	868	1632	3057	9	1	
1998	0.1	145	5	...	336	670	869	18	10	
1999	0.9	168	6	...	483	754	3919	19	4	

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [11]: T={'three_g':{'Yes':1,'No':0}}
test_df=test_df.replace(T)
print(test_df)
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	1	1043	1	1.8	1	14	0	5
1	2	841	1	0.5	1	4	1	61
2	3	1807	1	2.8	0	1	0	27
3	4	1546	0	0.5	1	18	1	25
4	5	1434	0	1.4	0	11	1	49
..
995	996	1700	1	1.9	0	0	1	54
996	997	609	0	1.8	1	0	0	13
997	998	1185	0	1.4	0	1	1	8
998	999	1533	1	0.5	1	0	0	50
999	1000	1270	1	0.5	0	4	1	35

	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h	sc_w
0	0.1	193	...	16	226	1412	3476	12	7
1	0.8	191	...	12	746	857	3895	6	0
2	0.9	186	...	4	1270	1366	2396	17	10
3	0.5	96	...	20	295	1752	3893	10	0
4	0.5	108	...	18	749	810	1773	15	8
..
995	0.5	170	...	17	644	913	2121	14	8
996	0.9	186	...	2	1152	1632	1933	8	1
997	0.5	80	...	12	477	825	1223	5	0
998	0.4	171	...	12	38	832	2509	15	11
999	0.1	140	...	19	457	608	2828	9	2

	talk_time	three_g	touch_screen	wifi
0	2	0	1	0
1	7	1	0	0
2	10	0	1	1
3	7	1	1	0
4	7	1	0	1
..
995	15	1	1	0
996	19	0	1	1
997	14	1	0	0
998	6	0	1	0
999	3	1	0	1

[1000 rows x 21 columns]

```
In [12]: x=train_df.drop('dual_sim',axis=1)
y=train_df['dual_sim']
```

```
In [13]: x=test_df.drop('dual_sim',axis=1)
y=test_df['dual_sim']
```

```
In [14]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

Out[14]: ((700, 20), (300, 20))

```
In [15]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[15]: RandomForestClassifier
RandomForestClassifier()
```

```
In [16]: rf=RandomForestClassifier()
```

```
In [17]: params={'max_depth':[2,3,5,10,20],
'min_samples_leaf':[5,10,20,50,100,200],
'n_estimators':[10,25,30,50,100,200]}
```

```
In [18]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
```

```
Out[18]:
GridSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier
```

```
In [19]: grid_search.best_score_
```

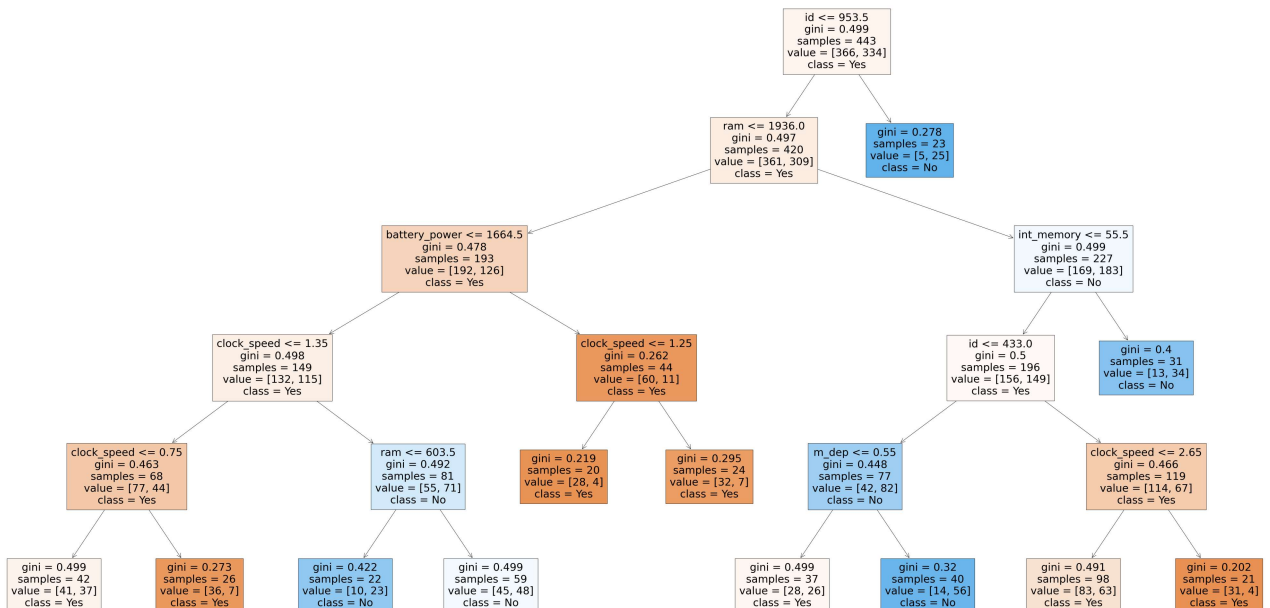
```
Out[19]: 0.54
```

```
In [20]: rf_best=grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=5, min_samples_leaf=20, n_estimators=30)
```

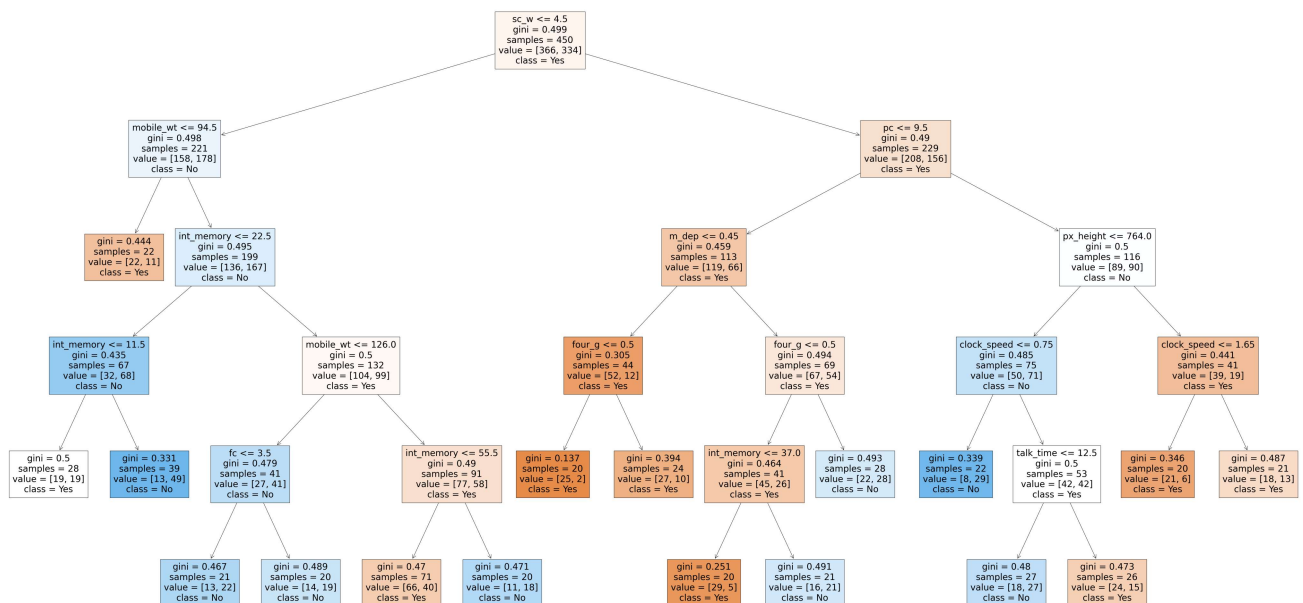
```
In [21]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[21]: [Text(0.6527777777777778, 0.9166666666666666, 'id <= 953.5\ngini = 0.499\nsamples = 443\nvalue = [366, 334]\nnclass = Yes'),
Text(0.5972222222222222, 0.75, 'ram <= 1936.0\ngini = 0.497\nsamples = 420\nvalue = [361, 309]\nnclass = Yes'),
Text(0.3611111111111111, 0.5833333333333334, 'battery_power <= 1664.5\ngini = 0.478\nsamples = 193\nvalue = [192, 126]\nnclass = Yes'),
Text(0.2222222222222222, 0.4166666666666667, 'clock_speed <= 1.35\ngini = 0.498\nsamples = 149\nvalue = [132, 115]\nnclass = Yes'),
Text(0.1111111111111111, 0.25, 'clock_speed <= 0.75\ngini = 0.463\nsamples = 68\nvalue = [77, 44]\nnclass = Yes'),
Text(0.05555555555555555, 0.08333333333333333, 'gini = 0.499\nsamples = 42\nvalue = [41, 37]\nnclass = Yes'),
Text(0.16666666666666666, 0.08333333333333333, 'gini = 0.273\nsamples = 26\nvalue = [36, 7]\nnclass = Yes'),
Text(0.3333333333333333, 0.25, 'ram <= 603.5\ngini = 0.492\nsamples = 81\nvalue = [55, 71]\nnclass = No'),
Text(0.2777777777777778, 0.08333333333333333, 'gini = 0.422\nsamples = 22\nvalue = [10, 23]\nnclass = No'),
Text(0.3888888888888889, 0.08333333333333333, 'gini = 0.499\nsamples = 59\nvalue = [45, 48]\nnclass = No'),
Text(0.5, 0.4166666666666667, 'clock_speed <= 1.25\ngini = 0.262\nsamples = 44\nvalue = [60, 11]\nnclass = Yes'),
Text(0.4444444444444444, 0.25, 'gini = 0.219\nsamples = 20\nvalue = [28, 4]\nnclass = Yes'),
Text(0.5555555555555556, 0.25, 'gini = 0.295\nsamples = 24\nvalue = [32, 7]\nnclass = Yes'),
Text(0.8333333333333334, 0.5833333333333334, 'int_memory <= 55.5\ngini = 0.499\nsamples = 227\nvalue = [169, 183]\nnclass = No'),
Text(0.7777777777777778, 0.4166666666666667, 'id <= 433.0\ngini = 0.5\nsamples = 196\nvalue = [156, 149]\nnclass = Yes'),
Text(0.6666666666666666, 0.25, 'm_dep <= 0.55\ngini = 0.448\nsamples = 77\nvalue = [42, 82]\nnclass = No'),
Text(0.6111111111111112, 0.08333333333333333, 'gini = 0.499\nsamples = 37\nvalue = [28, 26]\nnclass = Yes'),
Text(0.7222222222222222, 0.08333333333333333, 'gini = 0.32\nsamples = 40\nvalue = [14, 56]\nnclass = No'),
Text(0.8888888888888888, 0.25, 'clock_speed <= 2.65\ngini = 0.466\nsamples = 119\nvalue = [114, 67]\nnclass = Yes'),
Text(0.8333333333333334, 0.08333333333333333, 'gini = 0.491\nsamples = 98\nvalue = [83, 63]\nnclass = Yes'),
Text(0.9444444444444444, 0.08333333333333333, 'gini = 0.202\nsamples = 21\nvalue = [31, 4]\nnclass = Yes'),
Text(0.8888888888888888, 0.4166666666666667, 'gini = 0.4\nsamples = 31\nvalue = [13, 34]\nnclass = No'),
Text(0.7083333333333334, 0.75, 'gini = 0.278\nsamples = 23\nvalue = [5, 25]\nnclass = No')]
```



```
In [22]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[22]: [Text(0.41346153846153844, 0.9166666666666666, 'sc_w <= 4.5\ngini = 0.499\nsamples = 450\nvalue = [366, 334]\n\nclass = Yes'),
Text(0.1346153846153846, 0.75, 'mobile_wt <= 94.5\ngini = 0.498\nsamples = 221\nvalue = [158, 178]\n\nclass = No'),
Text(0.09615384615384616, 0.5833333333333333, 'gini = 0.444\nsamples = 22\nvalue = [22, 11]\n\nclass = Yes'),
Text(0.17307692307692307, 0.5833333333333333, 'int_memory <= 22.5\ngini = 0.495\nsamples = 199\nvalue = [136, 167]\n\nclass = No'),
Text(0.07692307692307693, 0.4166666666666667, 'int_memory <= 11.5\ngini = 0.435\nsamples = 67\nvalue = [32, 68]\n\nclass = No'),
Text(0.038461538461538464, 0.25, 'gini = 0.5\nsamples = 28\nvalue = [19, 19]\n\nclass = Yes'),
Text(0.11538461538461539, 0.25, 'gini = 0.331\nsamples = 39\nvalue = [13, 49]\n\nclass = No'),
Text(0.2692307692307692, 0.4166666666666667, 'mobile_wt <= 126.0\ngini = 0.5\nsamples = 132\nvalue = [104, 99]\n\nclass = Yes'),
Text(0.19230769230769232, 0.25, 'fc <= 3.5\ngini = 0.479\nsamples = 41\nvalue = [27, 41]\n\nclass = No'),
Text(0.15384615384615385, 0.08333333333333333, 'gini = 0.467\nsamples = 21\nvalue = [13, 22]\n\nclass = No'),
Text(0.23076923076923078, 0.08333333333333333, 'gini = 0.489\nsamples = 20\nvalue = [14, 19]\n\nclass = No'),
Text(0.34615384615384615, 0.25, 'int_memory <= 55.5\ngini = 0.49\nsamples = 91\nvalue = [77, 58]\n\nclass = Yes'),
Text(0.3076923076923077, 0.08333333333333333, 'gini = 0.47\nsamples = 71\nvalue = [66, 40]\n\nclass = Yes'),
Text(0.38461538461538464, 0.08333333333333333, 'gini = 0.471\nsamples = 20\nvalue = [11, 18]\n\nclass = No'),
Text(0.6923076923076923, 0.75, 'pc <= 9.5\ngini = 0.49\nsamples = 229\nvalue = [208, 156]\n\nclass = Yes'),
Text(0.5384615384615384, 0.5833333333333333, 'm_dep <= 0.45\ngini = 0.459\nsamples = 113\nvalue = [119, 66]\n\nclass = Yes'),
Text(0.46153846153846156, 0.4166666666666667, 'four_g <= 0.5\ngini = 0.305\nsamples = 44\nvalue = [52, 12]\n\nclass = Yes'),
Text(0.4230769230769231, 0.25, 'gini = 0.137\nsamples = 20\nvalue = [25, 2]\n\nclass = Yes'),
Text(0.5, 0.25, 'gini = 0.394\nsamples = 24\nvalue = [27, 10]\n\nclass = Yes'),
Text(0.6153846153846154, 0.4166666666666667, 'four_g <= 0.5\ngini = 0.494\nsamples = 69\nvalue = [67, 54]\n\nclass = Yes'),
Text(0.5769230769230769, 0.25, 'int_memory <= 37.0\ngini = 0.464\nsamples = 41\nvalue = [45, 26]\n\nclass = Yes'),
Text(0.5384615384615384, 0.08333333333333333, 'gini = 0.251\nsamples = 20\nvalue = [29, 5]\n\nclass = Yes'),
Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.491\nsamples = 21\nvalue = [16, 21]\n\nclass = No'),
Text(0.6538461538461539, 0.25, 'gini = 0.493\nsamples = 28\nvalue = [22, 28]\n\nclass = No'),
Text(0.8461538461538461, 0.5833333333333333, 'px_height <= 764.0\ngini = 0.5\nsamples = 116\nvalue = [89, 90]\n\nclass = No'),
Text(0.7692307692307693, 0.4166666666666667, 'clock_speed <= 0.75\ngini = 0.485\nsamples = 75\nvalue = [50, 71]\n\nclass = No'),
Text(0.7307692307692307, 0.25, 'gini = 0.339\nsamples = 22\nvalue = [8, 29]\n\nclass = No'),
Text(0.8076923076923077, 0.25, 'talk_time <= 12.5\ngini = 0.5\nsamples = 53\nvalue = [42, 42]\n\nclass = Yes'),
Text(0.7692307692307693, 0.08333333333333333, 'gini = 0.48\nsamples = 27\nvalue = [18, 27]\n\nclass = No'),
Text(0.8461538461538461, 0.08333333333333333, 'gini = 0.473\nsamples = 26\nvalue = [24, 15]\n\nclass = Yes'),
Text(0.9230769230769231, 0.4166666666666667, 'clock_speed <= 1.65\ngini = 0.441\nsamples = 41\nvalue = [39, 19]\n\nclass = Yes'),
Text(0.8846153846153846, 0.25, 'gini = 0.346\nsamples = 20\nvalue = [21, 6]\n\nclass = Yes'),
Text(0.9615384615384616, 0.25, 'gini = 0.487\nsamples = 21\nvalue = [18, 13]\n\nclass = Yes')]
```



```
In [23]: rf_best.feature_importances_
```

```
Out[23]: array([0.10586947, 0.07741492, 0.00400731, 0.07508846, 0.03478509,
0.00658804, 0.10753632, 0.04033311, 0.05944287, 0.02921685,
0.08363281, 0.0674476 , 0.06710616, 0.09544289, 0.04006483,
0.03878565, 0.03494972, 0.00490142, 0.02053677, 0.00684972])
```

```
In [24]: imp_df=pd.DataFrame({'Varname':x_train.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[24]:

	Varname	Imp
6	int_memory	0.107536
0	id	0.105869
13	ram	0.095443
10	pc	0.083633
1	battery_power	0.077415
3	clock_speed	0.075088
11	px_height	0.067448
12	px_width	0.067106
8	mobile_wt	0.059443
7	m_dep	0.040333
14	sc_h	0.040065
15	sc_w	0.038786
16	talk_time	0.034950
4	fc	0.034785
9	n_cores	0.029217
18	touch_screen	0.020537
19	wifi	0.006850
5	four_g	0.006588
17	three_g	0.004901
2	blue	0.004007

In []: